

Floating-point operations I

- The science of floating-point arithmetics
- IEEE standard
- Reference

What every computer scientist should know about floating-point arithmetic, ACM computing survey, 1991

Why learn more about floating-point operations I

Example:

- A one-variable problem

$$\begin{aligned} \min_x f(x) \\ x \geq 0 \end{aligned}$$

- In your program, should you set an upper bound of x ?
- x in your program may be wrongly increased to ∞

Why learn more about floating-point operations II

- What is the largest representable number in the computer ?
- Is there anything called infinity ?

Example:

- A ten-variable problem

$$\begin{aligned} &\min f(x) \\ &0 \leq x_i, i = 1, \dots, 10 \end{aligned}$$

Why learn more about floating-point operations III

- After the problem is solved, want to know how many are zeros ?
- Should you use
- People said: don't do floating-point comparisons

```
for (i=0; i < 10; i++)  
    if (x[i] == 0) count++ ;
```

```
epsilon = 1.0e-12 ;  
for (i=0; i < 10; i++)  
    if (x[i] <= epsilon) count++ ;
```

How do you choose ϵ ?

Why learn more about floating-point operations IV

- Is this true ?

Floating-point Formats I

- We know float (single): 4 bytes, double: 8 bytes
Why ?
- A floating-point system
base β , precision p , significand (mantissa) $d.d \dots d$
- Example

$$\begin{aligned} 0.1 &= \overbrace{1.00}^3 \times 10^{-1} & (\beta = 10, p = 3) \\ &\approx \underbrace{1.1001}_5 \times 2^{-4} & (\beta = 2, p = 5) \end{aligned}$$

exponent: -1 and -4

- Largest exponent e_{\max} , smallest e_{\min}

Floating-point Formats II

- β^p possible significands, $e_{\max} - e_{\min} + 1$ possible exponents

$$\lceil \log_2(e_{\max} - e_{\min} + 1) \rceil + \lceil \log_2(\beta^p) \rceil + 1$$

bits for storing a number

1 bit for \pm

- But the practical setting is **more complicated**
See the discussion of IEEE standard later
- Normalized: 1.00×10^{-1} (yes), 0.01×10^1 (no)
- Now ~~most used normalized representation~~
 \Rightarrow **cannot represent zero**

Floating-point Formats III

- A natural way for 0: $1.0 \times \beta^{e_{\min}-1}$
preserve the ordering

- Will use $p = 3, \beta = 10$ for most later explanation

the smallest
number
↓
 $1.0 \times \beta^{e_{\min}}$

Relative Errors and Ulp's I

$$\begin{array}{l} 3.14 \times 10^0 \\ 0.00159 \times 10^0 \end{array}$$

- When $\beta = 10, p = 3$, 3.14159 represented as 0.159 unit
 3.14×10^0

\Rightarrow error = $0.00159 = 0.159 \times 10^{-2}$, i.e. 0.159 units in the last place

10^{-2} : unit of the last place

- ulps: unit in the last place
- relative error $0.00159/3.14159 \approx 0.0005$
- For a number $d.d \dots d \times \beta^e$, the largest error is

$$\begin{array}{c} d.d \dots d \quad d.d \dots (d+1) \\ \hline \end{array}$$

$$0.\underbrace{0 \dots 0}_{p-1} \beta' \times \beta^e, \beta' = \beta/2$$

floating point number

Relative Errors and Ulp's II

- Error = $\frac{\beta}{2} \times \beta^{-p} \times \beta^e$ ex. $1 \times 10^5 \leq ? < 9 \times 10^5$

$$1 \times \underbrace{\beta^e} \leq \text{original value} < \beta \times \underbrace{\beta^e}$$

relative error between

$$\frac{\frac{\beta}{2} \times \beta^{-p} \times \cancel{\beta^e}}{\underbrace{\beta^e}} \quad \text{and} \quad \frac{\frac{\beta}{2} \times \beta^{-p} \times \beta^e}{\underbrace{\beta^{e+1}}}$$

so

$$\text{relative error} \leq \frac{\beta}{2} \beta^{-p} \quad (1)$$

- $\frac{\beta}{2} \beta^{-p} = \beta^{-p+1}/2$: machine epsilon

Relative Errors and Ulp's III

The bound in (1)

- When a number is rounded to the closest, relative error **bounded by ϵ**



ulps and ϵ

$$x = 12.35 \quad \tilde{x} = 12.4$$

- $p = 3, \beta = 10$

- Example: $x = 12.35 \Rightarrow \tilde{x} = 1.24 \times 10^1$

$$\text{error} = 0.05 = 0.005 \times 10^1$$

- $\text{ulps} = 0.01 \times 10^1, \epsilon = \frac{1}{2} 10^{-2} = 0.005$

- error 0.5 ulps $\rightarrow \frac{12.4}{12.35} \rightarrow \text{last place} = 0.1$

$$\text{relative error } 0.05/12.35 \approx 0.004 = 0.8\epsilon$$

$$= 0.01 \times 10^1 = \text{ulps}$$

- $8x = 98.8, 8\tilde{x} = 9.92 \times 10^1$

$$\text{error} = 4.0 \text{ ulps}$$

$$\text{relative error} = 0.4/98.8 = 0.8\epsilon.$$

- ulps and ϵ may be used interchangeably

Guard Digits I

- $p = 3, \beta = 10$
- Calculate $2.15 \times 10^{12} - 1.25 \times 10^{-5}$:
- Compute and then round

$$x = 2.15 \times 10^{12}$$

$$y = 0.0000000000000000000125 \times 10^{12}$$

$$x - y = 2.1499999999999999999875 \times 10^{12}$$

round to 2.15×10^{12}

Here we **assume** that computation is **exactly** done

Guard Digits II

- Round and then compute

$$x = 2.15 \times 10^{12}$$

$$y = 0.00 \times 10^{12}$$

$$x - y = 2.15 \times 10^{12}$$

Answer is the same

Reasonable as $x \approx x - y$

- Another example: $10.1 - 9.93 = 0.17$

ulps = $10^{-2} \times 10^{-1}$

$(.70 \times 10^{-1})$

2.00×10^{-1}

Guard Digits III

- Round and then compute

$$10.1 - 9.93 = 1.01 \times 10^1 - 0.99 \times 10^1 = 0.02 \times 10^1 \\ = 2.00 \times 10^{-1}$$

result

$$\text{error} = \underline{2.00 \times 10^{-1}} - \text{Correct } 0.17 = 0.03$$

$$\text{ulps} = 0.01 \times 10^{-1} = 10^{-3}$$

$$\text{error} = 0.03 = 30\text{ulps}$$

Handwritten notes:

\downarrow

$\downarrow \downarrow$

$\downarrow \downarrow$

$\text{ulps} = \beta^{-(p-1)} \cdot \beta^e$

Relative error

$$= 0.03/0.17 = 3/17$$

Guard Digits IV

The error is quite large

- Compute and round

$$10.1 - 9.93 = 0.17 = 1.7 \times 10^{-1}$$

error = 0

The problem: cannot compute and then round

Guard Digits V

$$S_{\text{small}} = \frac{0.00099}{0.90001} \rightarrow$$

- How big can the error be? (if round and then compute)

$$x = 1.000$$

$$y = 0.099 \boxed{99}$$

$$\Rightarrow x - y = 0.90001$$

$$\tilde{x} - \tilde{y} = 0.90100$$

Theorem

Using p digits with base β for $x - y$, the relative error can be as large as $\beta - 1$

Proof:

$$x = 1.0 \dots 0, y = 0. \underbrace{\eta \dots \eta}_{p \text{ digits}}, \eta = \beta - 1$$

$$\text{Correct solution } x - y = \beta^{-p}$$

$$\text{Computed solution} = 1.0 \dots 0 - 0. \underbrace{\eta \dots \eta}_{p-1 \text{ digits}} = \beta^{-p+1}$$

Guard Digits VI

Relative error

$$\frac{|\beta^{-p} - \beta^{-p+1}|}{\beta^{-p}} = \beta - 1$$

- Example: $p = 3$, $\beta = 10$

$$x = 1.00, y = 0.999, x - y = 0.001 = 10^{-3}$$

$$\begin{aligned}\text{Computed solution} &= 1.00 \times 10^0 - 0.99 \times 10^0 \\ &= 0.01 \times 10^0 = 0.01\end{aligned}$$

Relative error


$$\frac{|0.01 - 0.001|}{0.001} = 9$$

super large

Guard Digits VII

Such large errors occur if x and y are close

- Single guard digit



p increased by 1 in the device for addition and subtraction

round and then compute

$$\overset{p=4}{\underbrace{1.010}} \times 10^1 - \overset{p=4}{\underbrace{0.993}} \times 10^1 = 0.017 \times 10^1$$

Note $\overset{p=4}{\underbrace{0.017}} \times 10^1 = \overset{p=3}{\underbrace{1.70}} \times 10^{-1}$ can be stored as

- That is, one additional digit in the process of subtraction. All values are still stored using $p = 3$

Guard Digits VIII

- So in the device for subtraction, we should put additional digits
- Another example:

$$\begin{aligned} & 110 - 8.59 \\ &= 1.100 \times 10^2 - 0.085 \times 10^2 \\ &= 1.015 \times 10^2 \approx \underline{1.02} \times 10^2 \end{aligned}$$

$p=4$ \swarrow \searrow $p=3$
round

Correct answer 101.41

Relative error around 0.006

$$\begin{array}{r|l} 1.100 & \times 10^2 \\ - 0.0859 & \times 10^2 \\ \hline 1.0141 & \times 10^2 \end{array}$$

Guard Digits IX

$$\epsilon = \frac{1}{2}\beta^{-p+1} = \frac{1}{2}10^{-2} = 0.005$$

Theorem

Using $p + 1$ digits for $x - y \Rightarrow$ relative rounding error $< 2\epsilon$ (ϵ : machine epsilon)

Proof:

- Assume $x > y$
- Assume $x = x_0.x_1 \cdots x_{p-1} \times \beta^0$
The proof is similar if it's not β^0
- If $y = y_0.y_1 \cdots y_{p-1}$ no error

Guard Digits X

$$x_0, x_1, \dots, x_{p-1}$$

$$0. y_1 \dots y_{p-1} (y_p)$$

- If $y = 0.y_1 \dots y_p \Rightarrow$ 1 guard digit, exact $x - y$ rounded to a closest number \Rightarrow relative error $\leq \epsilon$

- **In general** $y = 0.0 \dots 0 y_{k+1} \dots y_{k+p}$
 \bar{y} : y truncated to $p + 1$ digits
guard length = p

$$|y - \bar{y}| < (\beta - 1)(\beta^{-p-1} + \beta^{-p-2} + \dots + \beta^{-p-k})$$

$$\leq \beta^{-p} \quad (2)$$

$-p - 1$: we have $p + 1$ digits now $0.00\dots y_{k+1}$ truncated

(Think about $p = 3, \beta = 10$, first digit truncated $\leq 9 \times 0.0001 = 9 \times 10^{-4}$)

$$< (\beta - 1)(\beta^{-p-1} + \dots + \beta^{-p-k})$$

$$= (\beta - 1) \cdot \beta^{-p-1} \cdot (1 + \dots + \beta^{-k})$$

Guard Digits XI

After y is truncated, we need to calculate

$$x - \bar{y}$$

It's rounded to

$$x - \bar{y} + \delta$$

$$|\delta| \leq (\beta/2)\beta^{-p} = \epsilon$$

The inequality comes from rounding a number of $p + 1$ digits

and round it to p digit


$$0.\underbrace{0 \dots 0}_{p \text{ digits}}(\beta/2) \dots$$

p digit

$$\frac{\beta}{2} \cdot \beta^{-p}$$

Guard Digits XII

max error

error: $(x - y) - (x - \bar{y} + \delta)$  $= \bar{y} - y - \delta$

Guard Digits XIII

case 1: if $x - y \geq 1$,

$$|\bar{y} - y - \delta| \leq |\bar{y} - y| + |\delta|$$

$$\begin{aligned} \epsilon &= \frac{\beta}{2} \beta^{-p} \\ 2\epsilon &= \frac{2\beta}{2} \beta^{-p} \\ \text{relative error} &= \frac{|\bar{y} - y - \delta|}{x - y} \leq \frac{|\bar{y} - y - \delta|}{1} \\ &\leq \beta^{-p} [(\beta - 1)(\beta^{-1} + \dots + \beta^{-k}) + \beta/2] \\ &= \beta^{-p} [(\beta - 1)\beta^{-k}(1 + \dots + \beta^{k-1}) + \beta/2] \\ &= \beta^{-p} [(\beta - 1)\beta^{-k} \frac{1 - \beta^k}{1 - \beta} + \beta/2] \\ &= \beta^{-p} [(1 - \beta^{-k}) + \beta/2] \\ &< \beta^{-p} (1 + \beta/2) \leq 2\epsilon \\ &= \frac{2+\beta}{2} < \frac{\beta+\beta}{2} \quad (\beta \geq 2) \end{aligned}$$

$$(\beta - 1) \beta^{-k} \frac{\beta^k - 1}{\beta - 1} = \beta^{-k} (\beta^k - 1)$$

Guard Digits XIV

$$x - y < 1 \rightarrow \text{or } \begin{matrix} x - \bar{y} \leq 1 \\ x - \bar{y} > 1 \end{matrix}$$

case 2: $x - \bar{y} \leq 1$: enough digits $\delta = 0$

the smallest $x - y$: (smallest x - largest y)

$$1.0 - 0.0 \dots 0 \underbrace{\rho \dots \rho}_k > (\beta - 1)(\beta^{-1} + \dots + \beta^{-k})$$

k zeros, p ρ 's, $\rho = \beta - 1$, from (2) the relative error

$$\begin{aligned} &\leq \frac{|\bar{y} - y - \delta|}{(\beta - 1)(\beta^{-1} + \dots + \beta^{-k})} \\ &< \frac{(\beta - 1)\beta^{-p}(\beta^{-1} + \dots + \beta^{-k})}{(\beta - 1)(\beta^{-1} + \dots + \beta^{-k})} = \beta^{-p} < 2\epsilon \end{aligned}$$

$\frac{2\beta}{\sum} \beta^{-p}$
 \parallel

case 3: $x - y < 1$ but $x - \bar{y} > 1$

We already prove : $|x-y| \geq 1$
We will then prove $|x-y| \leq 1$

We will then prove $|x - y| \leq 1$

$|x - y| < 1$ \rightarrow Case 2

$|x - y| > 1$ \rightarrow Case 3

Case 2: $(x-y) < 1$ and $x - \bar{y} \leq 1$

After $y \rightarrow \bar{y}$, $x - \bar{y} \Rightarrow \delta = 0$ p. 16, 17

$p+1$ digits

$\chi_0, \chi_1, \chi_2, \dots, \chi_{p-1} \in \mathbb{F}_p$

if $x - y = 1 \Rightarrow f = 0$

if $x - y < 1$

1.0000...
 ? . ? - - - ?
 p+1 digits

if $x - \bar{y} < 1 \Rightarrow$ we can normalize this number

So that $\oint = 0$

ex. $0.9999 \Rightarrow 1.999$

Guard Digits XV

We show that this situation is impossible ✓

If $x - \bar{y} = 1.\underbrace{0 \dots 0}_p \Rightarrow x - y \geq 1$: a contradiction

Why $x - y$ must be ≥ 1 : $\begin{array}{r} 1.00\dots1 \\ -0.00\dots1 \\ \hline 1.0000\dots \end{array}$

$$-\beta^{-p} < y - \bar{y} < \beta^{-p} \quad |y - \bar{y}| < \beta^{-p} \neq 0.0\dots1$$

$$\Rightarrow -\bar{y} - \beta^{-p} < -y \Rightarrow \boxed{x - \bar{y} - \beta^{-p}} < x - y$$

The difference between y and \bar{y} is $0.0\dots1$

- Conclusion: adding some guard digits can reduce the error

Especially when subtracting two nearby numbers

- Cost: the adder is one bit wider (cheap)

Most modern computers have guard digits

Cancellation I

- Catastrophic cancellation and benign cancellation
- Catastrophic cancellation :

$$b = 3.34, a = 1.22, c = 2.28, b^2 - 4ac = 0.0292$$

$$b^2 \approx 11.2, 4ac \approx 11.1 \Rightarrow \text{answer} = 0.1$$

$$\text{error} = 0.1 - 0.0292 = 0.0708$$

$$\text{answer} = 0.0292 = 2.92 \times 10^{-2}$$

$$\text{ulps} = 0.01 \times 10^{-2} = 10^{-4}$$

$$0.0708 \approx 708 \text{ ulps}$$

- Happens when subtracting two close numbers

Cancellation II

- Benign cancellation: subtracting **exactly** known numbers, by guard digits
- \Rightarrow small relative error
- In the example, b^2 and $4ac$ already contain errors

Avoid Catastrophic Cancellation I

- By rearranging the formula
- Example

$$\frac{-b + \sqrt{b^2 - 4ac}}{2a} \quad (3)$$

- If $b^2 \gg 4ac \Rightarrow$ no cancellation when calculating $b^2 - 4ac$ and $\sqrt{b^2 - 4ac} \approx |b|$
Then $-b + \sqrt{b^2 - 4ac}$ has a catastrophic cancellation if $b > 0$

Avoid Catastrophic Cancellation II

- Multiplying $-b - \sqrt{b^2 - 4ac}$, if $b > 0$

$$\frac{2c}{-b - \sqrt{b^2 - 4ac}} \quad (4)$$

- Use (3) if $b < 0$, (4) if $b > 0$
- Difficult to remove all catastrophic cancellations, but possible to remove most by reformulations
- Another example: $x^2 - y^2$

Assume $x \approx y$

$(x - y)(x + y)$ is better than $x^2 - y^2$

Avoid Catastrophic Cancellation III

x^2, y^2 may be rounded $\Rightarrow x^2 - y^2$ may be a catastrophic cancellation

$x - y$ by guard digit

- A catastrophic cancellation is replaced by a benign cancellation

Of course x, y may have been rounded and $x - y$ is still a catastrophic cancellation.

Again, difficult to remove all catastrophic cancellations, but possible to remove some

Avoid Catastrophic Cancellation IV

- Calculating area of a triangle

$$A = \sqrt{s(s-a)(s-b)(s-c)}, s = \frac{a+b+c}{2} \quad (5)$$

a, b, c : length of three edges

If $a \approx b + c$, $s = (a + b + c)/2 \approx a$, $s - a$ may have a catastrophic error

Example: $a = 9.00$, $b = c = 4.53$

$s = 9.03$, $A = 2.342$

Computed solution: $A = 3.04$, error ≈ 0.7

ulps = 0.01, error = 70 ulps

Avoid Catastrophic Cancellation V

- A new formulation by Kahan [1986] , $a \geq b \geq c$

$A =$

$$\frac{\sqrt{(a + (b + c))(c - (a - b))(c + (a - b))(a + (b - c))}}{4} \quad (6)$$

$A \approx 2.35$, close to 2.342

- HW 1-1: Calculate $A = 3.04$ using (5) and $A = 2.35$ using (6)

Avoid Catastrophic Cancellation VI

Note: to get $A = 3.04$ you need to calculate s by

$$s = \frac{a + (b + c)}{2}$$

Note that for multiplication and square root we assume that exact calculation can be done and results are rounded.

- Conclusion: sometimes a formula can be rewritten to have higher accuracy using benign cancellation
- Only works if guard digit is used; most computers use guard digits now

Avoid Catastrophic Cancellation VII

- But reformulation is difficult!!
You may think that you will never need to do this
- Two real cases:
- Line 213-216 of tron.cpp of LIBLINEAR version 2.11
<http://www.csie.ntu.edu.tw/~cjlin/liblinear/oldfiles>
HW1-2: Check Eq. (13) of the paper
<http://www.csie.ntu.edu.tw/~cjlin/papers/logistic.pdf>
and explain how we avoid catastrophic cancellations

Avoid Catastrophic Cancellation VIII

We do not consider the latest version of LIBLINEAR because some more complicated settings have been used

- Probability outputs of LIBSVM

HW1-3: Repeat the experiment on page 5, line 12 of the paper

<http://www.csie.ntu.edu.tw/~cjlin/papers/plattprob.pdf>

Discuss what you found

Exactly Rounded Operations I

- Round then calculate \Rightarrow may not be very accurate
- **Exactly rounded**: compute exactly then rounded to the nearest \Rightarrow usually more accurate
- The definition of **rounding**
- $12.5 \Rightarrow 12$ or 13 ?
- Rounding up: 0, 1, 2, 3, 4 \Rightarrow down, 5, 6, 7, 8, 9 \Rightarrow up

Why called “rounding up”? Always up for 5

- Rounding even:
the closest value with even least significant digit

Exactly Rounded Operations II

50% probability up, 50% down

example: $12.5 \Rightarrow 12$; $11.5 \Rightarrow 12$

- Reiser and Knuth [1975] show rounding even may be better

Theorem

Let $x_0 = x$, $x_1 = (x_0 \ominus y) \oplus y$, \dots , $x_n = (x_{n-1} \ominus y) \oplus y$, if \oplus and \ominus are exactly rounded using rounding even, then $x_n = x, \forall n$ or $x_n = x_1, \forall n \geq 1$.

$x \ominus y$: computed solution

- Consider rounding up,

Exactly Rounded Operations III

$$\beta = 10, p = 3, x = 1.00, y = -0.555$$

$$x - y = 1.555, x \ominus y = 1.56, (x \ominus y) + y = 1.56 - 0.555 = 1.005, x_1 = (x \ominus y) \oplus y = 1.01$$

$$x_1 - y = 1.565, x_1 \ominus y = 1.57, (x_1 \ominus y) + y = 1.57 - 0.555 = 1.015, x_2 = (x_1 \ominus y) \oplus y = 1.02$$

Increased by 0.01 until $x_n = 9.45$

- Rounding even:

$$x - y = 1.555, x \ominus y = 1.56, (x \ominus y) + y = 1.56 - 0.555 = 1.005, x_1 = (x \ominus y) \oplus y = 1.00$$

$$x_1 - y = 1.555, x_1 \ominus y = 1.56, (x_1 \ominus y) + y = 1.56 - 0.555 = 1.005, x_2 = (x_1 \ominus y) \oplus y = 1.00$$

Exactly Rounded Operations IV

- How to implement “exactly rounded operations”?
We can use an array of words or floating-points
But you don’t have an infinite amount of spaces
- Goldberg [1990] showed that using **3 guard digits** the result is the same as using exactly rounded operations

IEEE standard I

- IEEE 754 during 80s, now standard everywhere
- Two IEEE standards:
 - 754: specify $\beta = 2, p = 24$ for single, $\beta = 2, p = 53$ for double
 - 854 ($\beta = 2$ or 10 , does not specify how floating-point numbers are encoded into bits)
- Why IEEE 854 allows $\beta = 2$ or 10 but not other numbers:
 - 10 is the base we use
 - smaller β causes smaller relative error

IEEE standard II

smaller β : more precision. For example,

$$\beta = 16, p = 1 \text{ versus } \beta = 2, p = 4$$

- 4 bits for significand

$$\epsilon = \frac{16}{2}16^{-1} = 1/2, \epsilon = \frac{2}{2}2^{-4} = 1/16$$

We can see that ϵ of $\beta = 2, p = 4$ is smaller

- However, IBM/370 uses $\beta = 16$. Why? Two possible reasons:

First,

IEEE standard III

a number: 4 bytes = 32 bits

$\beta = 16, p = 6$, significand: $4 \times 6 = 24$ bits,

exponents: $32 - 24 - 1 = 7$ bits (1 bit for sign),
 16^{-2^6} to $16^{2^6} = 2^{2^8}$

for $\beta = 2 \Rightarrow 9$ bits (-2^8 to $2^8 = 2^9$) for exponents,
 $\Rightarrow 32 - 9 - 1 = 22$ for significand

Same exponents, less significand for $\beta = 2$ (24 vs. 22)

Second,

Shifting: $\beta = 16$, less frequently to adjust exponents when adding or subtracting two numbers

IEEE standard IV

For modern computers, this saving is not important

- Single precision: $\beta = 2$, $p = 24$ (23 bits as normalized), exponent 8, 1 bit for sign
($32 = 23 + 8 + 1$)
- An example: $176.625 = 1.0101100101 \times 2^7$

0 10000110 010110010100000000000000

1 of $1.\dots$ is not stored (normalized)

- Biased exponent (described later in detail)
 $10000110 = 128 + 4 + 2 = 134, 134 - 127 = 7$

Note that we have **negative** exponent

IEEE standard V

- Use rounding even

Binary	rounded	reason
10.00011	10.00	($< 1/2$, down)
10.00110	10.01	($> 1/2$, up)
10.11100	11.00	($1/2$, up)
10.10100	10.10	($1/2$, down)

This example is from <http://www.cs.cmu.edu/afs/cs/academic/class/15213-s12/www/lectures/04-float-4up.pdf>

- A summary

IEEE standard VI

IEEE	Fortran	C	Bits	Exp.	Mantissa
Single	REAL*4	float	32	8	24
Single-extended			44	≤ 11	32
Double	REAL*8	double	64	11	53
Double-extended	REAL*10	long double	≥ 80	≥ 15	≥ 64

$$32 = 8 + 24 \text{ but } 44 \neq 11 + 32$$

- $44 \neq 11 + 32$:

Hardware implementation of extended precision normal don't use a hidden bit

(Remember we normalized each number so 1 is not stored)

IEEE standard VII

- It seems everyone is using double now
But single is still needed sometime (if memory is not enough)
- Minimal normalized positive number

$$1 \times 2^{-126} \approx 1.17 \times 10^{-38}$$

$$e_{\min} = -126$$

- 8 bits for exponent: 0 to 255. IEEE uses biased approach exponent

$$(0 \text{ to } 255) - 127 = -127 \text{ to } 128$$

IEEE standard VIII

- Why $e_{\min} = -126$ but $e_{\max} = 127$?
reasons: $1/2^{e_{\min}}$ not overflow, $1/2^{e_{\max}}$ underflow, but less serious
- Thus, -127 for 0 and denormalized numbers (discussed later), -126 to 127 for exponents, 128 for special quantity
- Motivation for extended precision: from calculator, display 10 digits but 13 internally
Some operations benefit from using more digits internally

IEEE standard IX

Example: binary-decimal conversion (Details not discussed here)

- Operations: IEEE standard requires results of addition, subtraction, multiplication and division exactly rounded.
- Exactly rounded: an array of words or floating-point numbers, expensive
- Goldberg [1990] showed using 3 guard digits the result is the same as using exactly rounded operations

Only little more cost

IEEE standard X

- Reasons to specify operations
run on different machines \Rightarrow results the same
- HW 2-1: write the binary format of -250 as a double floating-point number
- IEEE: square root, remainder, conversion between integer and floating-point, internal formats and decimal are correctly rounded (i.e. exactly rounded operations)
- Binary to decimal conversion
Think about reading numbers from files

IEEE standard XI

When writing a binary number to a decimal number and read it back, can we get the same binary number?

- Writing 9 digits is enough for short
Though $10^8 > 2^{24}$, 8 digits are not enough
- 17 for double precision (proof not provided).

Example:

numbers in a data set from Matrix market:

IEEE standard XII

```
> tail s1rmq4m1.dat
  8.2511736085618438E+01    2.5134528659924950
 -6.0042951255041466E+00    8.6599442206615524
  1.0026197619563723E+01   -1.3136837661844502
 -1.5108331040361231E+01    5.1423173996955084
 -1.1690286345961363E+03    1.6250726655807816
  8.2511736074473220E+01    1.5108331040361227
```

- Matrix market:

<http://math.nist.gov/MatrixMarket/>

A collection of matrix data

IEEE standard XIII

- Transcendental numbers:
e.g., \exp , \log
- IEEE does not require transcendental functions to be exactly rounded

Cannot specify the precision because they are arbitrarily long

Special quantities I

- On some computers (e.g., IBM 370) every bit pattern is a valid floating-point number
- For IBM 370, $\sqrt{-4} = 2$ and it prints an error message

IEEE : NaN, not a number

why $\sqrt{-4} = 2$ on IBM 370 \Rightarrow every pattern is a number

- Special value of IEEE:
 $+0$, -0 , denormalized numbers, $+\infty$, $-\infty$, NaNs
(more than one NaN)

Special quantities II

- A summary

Exponent	significand	represents
$e = e_{\min} - 1$	$f = 0$	$+0, -0$
$e = e_{\min} - 1$	$f \neq 0$	$0.f \times 2^{e_{\min}}$
$e_{\min} \leq e \leq e_{\max}$		$1.f \times 2^e$
$e = e_{\max} + 1$	$f = 0$	$\pm\infty$
$e = e_{\max} + 1$	$f \neq 0$	NaN

- Why IEEE has NaN

Sometimes even $0/0$ occurs, the program can continue

Special quantities III

- Example: find $f(x) = 0$, try different x 's, even $0/0$ happens, other values may be ok.
- If $b^2 - 4ac < 0$

$$\frac{-b + \sqrt{b^2 - 4ac}}{2a}$$

returns NaN

$-b +$ NaN should be NaN

In general when a NaN is in an operation, result is NaN

- Examples producing NaN:

Special quantities IV

Operation	NaN by
+	$\infty + (-\infty)$
\times	$0 \times \infty$
/	$0/0, \infty/\infty$
REM	$x \text{ REM } 0, \infty \text{ REM } y$
$\sqrt{}$	\sqrt{x} when $x < 0$

Infinity I

- $\beta = 10, p = 3, e_{\max} = 98, x = 3 \times 10^{70},$
 x^2 overflow and replaced by $9.99 \times 10^{98}??$

In IEEE, the result is ∞

- Note $0/0 = \text{NaN}, 1/0 = \infty, -1/0 = -\infty$
 \Rightarrow nonzero divided by 0 is ∞ or $-\infty$

Similarly, $-10/0 = -\infty$, and $-10/-0 = +\infty$
(± 0 will be explained later)

- $3/\infty = 0, 4 - \infty = -\infty, \sqrt{\infty} = \infty$
- How to know the result?

replace ∞ with x , let $x \rightarrow \infty$

Infinity II

Example:

$$3/\infty : \lim_{x \rightarrow \infty} 3/x = 0$$

If limit does not exist \Rightarrow NaN

- $x/(x^2 + 1)$ vs $1/(x + x^{-1})$

$x/(x^2 + 1)$: if x is large, x^2 overflow, $x/\infty = 0$ but not $1/x$.

$1/(x + x^{-1})$: x large, $1/x$ ok

$1/(x + x^{-1})$ looks better but what about $x = 0$?

$$x = 0, 1/(0 + 0^{-1}) = 1/(0 + \infty) = 1/\infty = 0$$

- If no infinity arithmetic, an extra instruction needed to test if $x = 0$, may interrupt the pipeline

Signed zero I

- Why do we have $+0$ and -0 ?

First, it is available (1 bit for sign)

if no sign, $1/(1/x) = x$ fails when $x = \pm\infty$

$$x = \infty, 1/x = 0, 1/0 = +\infty$$

$$x = -\infty, 1/x = 0, 1/0 = +\infty$$

- Compare $+0$ and -0 : if $(x == 0)$

IEEE defines $+0 = -0$

- IEEE: $3 \times (+0) = +0$, $+0/(-3) = -0$

Signed zero II

- ± 0 useful in the following situations:

$$\log x \equiv \begin{cases} -\infty & x = 0 \\ \text{NaN} & x < 0 \end{cases}$$

A small underflow negative number $\Rightarrow \log x$ should be NaN

x underflow \Rightarrow round to 0, if no sign, $\log x$ is $-\infty$ but not NaN

Signed zero III

- With ± 0 , we have

$$\log x = \begin{cases} -\infty & x = +0 \\ \text{NaN} & x = -0 \\ \text{NaN} & x < 0 \end{cases}$$

Positive underflow \Rightarrow round to $+0$

- Very useful in complex arithmetic

$$\sqrt{1/z} \text{ and } 1/\sqrt{z}$$

$$z = -1, \sqrt{1/-1} = \sqrt{-1} = i, 1/\sqrt{-1} = 1/i = -i$$

$$\Rightarrow \sqrt{1/z} \neq 1/\sqrt{z}$$

Signed zero IV

- This happens because square root is multi-valued.

$$i^2 = (-i)^2 = -1$$

- However, by some restrictions (or ways of calculation), they can be equal

- $z = -1 = -1 + 0i$,

$$1/z = 1/(-1 + 0i) = -1 + (-0)i$$

$$\text{so } \sqrt{1/z} = \sqrt{-1 + (-0)i} = -i$$

$\Rightarrow -0$ is useful

- Disadvantage of $+0$ and -0 :

$$x = y \Leftrightarrow 1/x = 1/y \text{ is destroyed}$$

$$x = 0, y = -0 \Rightarrow x = y \text{ under IEEE}$$

Signed zero V

$$1/x = +\infty, 1/y = -\infty, +\infty \neq -\infty$$

- There are always pros and cons for floating-point design

HW 2-2 I

- If

$\text{if } (a < 0)$

always holds and b is neither too large nor too small, how do we guarantee

$\text{if } a/\max(b, 0.0) < 0$

always holds

- If $\max(b, 0.0)$ returns -0.0 , then it may not hold
- For the max function, should we use

$(x > y)? \quad x : y$

or

HW 2-2 II

$(x < y) ? y : x$

- Your max need to return $+0.0$ but not -0.0
- How to specifically assign $+0.0$ and -0.0 ?
- How to use subroutines to get the sign of a number?
- In a regular program, if you write 0.0 , is it $+0.0$ or -0.0 ?

Find the statement in the manual saying that 0.0 means $+0.0$

- Do some experiments to check your arguments
- Use Java but not other systems

Denormalized number I

- $\beta = 10, p = 3, e_{\min} = -98, x = 6.87 \times 10^{-97}, y = 6.81 \times 10^{-97}$
- x, y are ok but $x - y = 0.6 \times 10^{-98}$ rounded to 0, even though $x \neq y$
- How important to preserve

$$x = y \Leftrightarrow x - y = 0$$

- if $(x \neq y) \{z = 1/(x-y); \}$

The statement is true, but z becomes ∞

Tracking such bugs is frustrating

Denormalized number II

- IEEE uses denormalized numbers

Guarantee $x = y \Leftrightarrow x - y = 0$

Details of how this is done are not discussed here

- **Most controversial part** in IEEE standard

It caused long delay of the standard

- If denormalized number is used, 0.6×10^{-98} is also a floating-point number
- Remember we do not store 1 of $1.d \cdots d$
- How to represent denormalized numbers ?

Recall for valid value, $e \geq e_{\min}$ and we have

$$1.d \cdots d \times 2^e$$

Denormalized number III

- For denormalized numbers, we let $e = e_{\min} - 1$ and the corresponding value be

$$0.d \cdots d \times 2^{e+1} = 0.d \cdots d \times 2^{e_{\min}}$$

- Why not

$$1.d \cdots d \times 2^{e_{\min}-1}$$

can't represent

$$0.0x \cdots x \times 2^{e_{\min}}$$

- $6.87 \times 10^{-97} - 6.81 \times 10^{-97} \Rightarrow$ underflow due to cancellation

Denormalized number IV

Underflow: smaller than the smallest floating-point number

- An example of using denormalized numbers

$$\begin{aligned}\frac{a + bi}{c + di} &= \frac{(a + bi)(c - di)}{(c + di)(c - di)} \\ &= \frac{ac + bd}{c^2 + d^2} + \frac{bc - ad}{c^2 + d^2}i\end{aligned}$$

If c or $d > \sqrt{\beta}\beta^{e_{\max}/2} \Rightarrow$ overflow

Denormalized number V

overflow: larger than the maximal floating-point number

- Smith's formula

$$\frac{a + bi}{c + di} = \begin{cases} \frac{a+b(d/c)}{c+d(d/c)} + \frac{b-a(d/c)}{c+d(d/c)}i & \text{if } (|d| < |c|) \\ \frac{b+a(c/d)}{d+c(c/d)} + \frac{-a+b(c/d)}{d+c(c/d)}i & \text{if } (|d| \geq |c|) \end{cases}$$

avoid overflow

- **However**, using Smith's formula, without denormalized numbers

Denormalized number VI

If

$$a = 2 \times 10^{-98}, b = 1 \times 10^{-98}, c = 4 \times 10^{-98}, \\ d = 2 \times 10^{-98}$$

then

$$d/c = 0.5, c + d(d/c) = 5 \times 10^{-98}, \\ b(d/c) = 1 \times 10^{-98} \times 0.5 = 0 \\ a + b(d/c) = 2 \times 10^{-98}$$

Solution = 0.4 , wrong

Denormalized number VII

If denormalized numbers are used, 0.5×10^{-98} can be stored,

$$a + b(d/c) = 2.5 \times 10^{-98} \Rightarrow 0.5$$

the correct answer

- Usually hardware does not support denormalized numbers directly

Using software to simulate

- Programs may be slow if a lot of underflow

Exception, Flags, Trap handlers I

- We have mentioned things like overflow, underflow
What are other exceptional situations ?
- Motivation: usually when exceptional condition like $1/0$ happens, you may want to know
- IEEE requires vendors to provide a way to get status flags
- IEEE defines five exceptions: overflow, underflow, division by zero, invalid operation, inexact
- overflow: larger than the maximal floating-point number

Exception, Flags, Trap handlers II

Underflow: smaller than the smallest floating-point number

- Invalid:

$\infty + (-\infty)$, $0 \times \infty$, $0/0$, ∞/∞ ,
 $x \text{ REM } 0$, $\infty \text{ REM } y$, \sqrt{x} , $x < 0$, any comparison
involves a NaN

- Invalid returns NaN; NaN may not be from invalid operations

- Inexact: the result is not exact

$\beta = 10$, $p = 3$, $3.5 \times 4.2 = 14.7$ exact,
 $3.5 \times 4.3 = 15.05 \Rightarrow 15.0$ not exact

Exception, Flags, Trap handlers III

inexact exception is raised so often, usually we ignore it

Exception	when trap disabled	argument to handler
overflow	$\pm\infty$ or $\pm 1.1 \dots 1 \times 2^{e_{\max}}$	$\text{round}(x2^{-\alpha})$
underflow	$0, \pm 2^{e_{\min}}$, or denormal	$\text{round}(x2^{\alpha})$
division by zero	∞	operands
invalid	NaN	operands
inexact	$\text{round}(x)$	$\text{round}(x)$

- Trap handler: special subroutines to handle exceptions

Exception, Flags, Trap handlers IV

You can design your own trap handlers

- In the above table, “when trap disabled” means results of operations if trap handlers not used
- $\alpha = 192$ for single, $\alpha = 1536$ for double
reason: you cannot really store x
- Examples of using trap handlers described later

Compiler Options I

- Compiler may provide a way so the program stops if an exception occurs
- Easy for debugging
- Example: SUN's C compiler (I learned this on an old machine)
- Reason: gcc doesn't have this to explicitly detect exceptions
- **-ftrap=t**

Compiler Options II

- t: %all, %none, common, [no%]invalid, [no%]overflow, [no%]underflow, [no%]division, [no%]inexact.
- common: invalid, division by zero, and overflow.
- The default is -ftrap=%none.
- Example: -ftrap=%all,no%inexact means set all traps, except inexact.
- If you compile one routine with -ftrap=t, compile all routines of the program with the same -ftrap=t option
otherwise, you can get unexpected results.

Compiler Options III

- Example: on the screen you will see

Note: IEEE floating-point exception flags raised:
Inexact; Underflow;
See the Numerical Computation Guide, `ieee_flags`.

- gcc:
- `-fno-trapping-math`: default `-ftrapping-math`
Setting this option may allow faster code if one relies on “non-stop” IEEE arithmetic
- `-ftrapv`

Compiler Options IV

Generates traps for signed overflow on addition, subtraction, multiplication

Trap Handler I

- Example:

```
do {  
    ....  
} while {not x >= 100;}
```

If $x = \text{NaN}$, an infinite loop

Any comparison involving NaN is wrong

- A trap handler can be installed to abort it
- Example:

Trap Handler II

Calculate $x_1 \times \cdots \times x_n$ may overflow in the middle (the total may be ok!):

```
for (i = 1; i <= n; i++)  
    p = p * x[i] ;
```

- $x_1 \times \cdots \times x_r, r \leq n$ overflow but $x_1 \times \cdots \times x_n$ may be in the range
- $e^{\sum \log(x_i)} \Rightarrow$ a solution but less accurate and costs more
- A possible solution

Trap Handler III

```
for (i = 1; i <= n; i++) {  
    if (p * x[i] overflow) {  
        p = p * pow(10,-a);  
        count = count + 1 ;  
    }  
    p = p * x[i] ;  
}  
p = p * pow(10, a*count) ;
```

An Example of Handlers I

- Example using SUN's numerical computation guide
Again, old. Reason of not using existing systems
such a glibc: so you can have HW
- standard math library libm.a
exp, pow, log, ...
- On SUN machines, there are additional math
library: libsunmath.a
exp2, exp10, ..., ieee_flags, ieee_handler,
ieee_retrospective
- A program:

An Example of Handlers II

```
#include <stdio.h>
#include <sys/ieee754.h>
#include <sunmath.h>
#include <siginfo.h>
#include <ucontext.h>

void handler(int sig, siginfo_t *sip,
             ucontext_t *uap)
{
    unsigned    code, addr;

    code = sip->si_code;
```


An Example of Handlers III

```
    addr = (unsigned) sip->si_addr;
    fprintf(stderr, "fp exception %x at
               address %x \n", code, addr);
}
int main()
{
    double  x;

    /* trap on common floating point
       exceptions */
    if (ieee_handler("set", "common", handler)
        != 0)
```

An Example of Handlers IV

```
printf("Did not set exception  
handler \n");
```

```
/* cause an underflow exception (not  
reported) */  
x = min_normal();  
printf("min_normal = %g \n", x);  
x = x / 13.0;  
printf("min_normal / 13.0 = %g \n", x);  
  
/* cause an overflow exception  
(reported) */
```

An Example of Handlers V

```
x = max_normal();  
printf("max_normal = %g \n", x);  
x = x * x;  
printf("max_normal * max_normal = %g \n",  
      x);  
  
ieee_retrospective(stderr);  
return 0;  
}
```

- Result:

An Example of Handlers VI

```
min_normal = 2.22507e-308
```

```
min_normal / 13.0 = 1.7116e-309
```

```
max_normal = 1.79769e+308
```

```
fp exception 4 at address 10d0c
```

```
max_normal * max_normal = 1.79769e+308
```

Note: IEEE floating-point exception flags raised:

Inexact; Underflow;

IEEE floating-point exception traps enabled:

overflow; division by zero; invalid operation

See the Numerical Computation Guide, `ieee_flags`

`ieee_handler(3M)`

An Example of Handlers VII

- invalid, division, and overflow sometimes called common exceptions here
ieee_handler("set", "common", handler) means handlers used for common exceptions
- min_normal / 13.0: using denormalized numbers
handler: subroutines to handle exceptions
- HW 3-1: regenerate this example using GNU C library

An Example of Handlers VIII

- How to find GNU C library information: on linux, type
% info libc
check the category of “Arithmetics” and “Signal Handling”

The Use of Flags: An Example I

- Calculate x^n , n : integer

```
double pow(double x, int n)
{
    double tmp = x, ret = 1.0;

    for(int t=n; t>0; t/=2)
    {
        if(t%2==1) ret*=tmp;
        tmp = tmp * tmp;
    }
    return ret;
```

The Use of Flags: An Example II

}

$x^{16} = (x^2)^8 = \dots$, $x^{15} = x(x^2)^7$, treat x^2 as the new x

$$x^{15} = x(x^2)^7 = x(x^2)(x^4)^3 = x(x^2)(x^4)(x^8)^1$$

- If $n < 0$, we need to use

$$x^n = (1/x)^{-n} = 1/(x)^{-n}$$

$\text{pow}(1/x, -n)$ less accurate, $1/\text{pow}(x, -n)$ is better

There is already error on $1/x$

The Use of Flags: An Example III

Example: $2^{-5} = (1/2)^5$ and $1/(2^5)$

- A small problem on using $1/\text{pow}(x, -n)$:
if $\text{pow}(x, -n)$ underflow (i.e. when $x < 1, n < 0$),
either underflow trap handler or underflow status
flag set \Rightarrow incorrect

x^{-n} underflow, x^n overflow or be in range
($e_{\min} = -126, 2^{-e_{\min}} = 2^{126} < 2^{127} = 2^{e_{\max}}$)

- Turn off overflow & underflow trap enable bits, save
overflow & underflow status bits

Compute $1/\text{pow}(x, -n)$

The Use of Flags: An Example IV

If neither overflow nor underflow status is set \Rightarrow restore them

If one is set, restore & calculate $\text{pow}(1/x, -n)$, which causes correct exception to occur

- Practically the calculation of $\text{pow}()$ is more complicated
e.g. google `e_pow.c` and `e_log.c`
- In `glibc-2.17/sysdeps/ieee754/dbl-64`, `e_pow.c` has 420 lines

The Use of Flags: An Example V

- Another example: calculate arccos using arctan

$$\arccos x = 2 \arctan \sqrt{\frac{1-x}{1+x}}$$

$$\cos \theta = x = 2 \cos^2 \frac{\theta}{2} - 1 = 1 - 2 \sin^2 \frac{\theta}{2}$$

$$\cos \frac{\theta}{2} = \sqrt{\frac{x+1}{2}}, \sin \frac{\theta}{2} = \sqrt{\frac{1-x}{2}}, \tan \frac{\theta}{2} = \sqrt{\frac{1-x}{1+x}}$$

Hence

$$\arccos x = 2 \arctan \sqrt{\frac{1-x}{1+x}}$$

The Use of Flags: An Example VI

- Consider $x = -1$

$$\arctan(\infty) = \pi/2 \Rightarrow \arccos(-1) = \pi$$

- A small problem:

$\frac{1-x}{1+x}$ causes the divide-by-zero flag set though $\arccos(-1)$ not exceptional

- Solution: save divide-by-zero flag, restore after \arccos computation

A Real Study I

- Let's start with a simple example

```
#include <stdio.h>
```

```
int main()
```

```
{
```

```
float a = 123.123;
```

```
printf("%.10f\n", a);
```

```
printf("%.10f\n", a*a);
```

```
a = 123.125;
```

```
printf("%.10f\n", a);
```

A Real Study II

```
printf("%.10f\n", a*a);
```

```
}
```

- Results are

```
$gcc test.c; ./a.out
```

```
123.1230010986
```

```
15159.2734375000
```

```
123.1250000000
```

```
15159.7656250000
```

```
$gcc -m32 test.c; ./a.out
```

```
123.1230010986
```

A Real Study III

15159.2733995339

123.1250000000

15159.7656250000

- -m 32 generates code for a 32-bit environment (because we don't have a 32-bit machine)
- That is, same code gives different results under 32 and 64-bit environments
- Why?

A Real Study IV

- On 32 bit, 387 floating-point coprocessor is used. From gcc manual, “The temporary results are computed in 80-bit precision instead of the precision specified by the type, resulting in slightly different results compared to most of other chips.”
- In other words, they somehow **violate** IEEE standard
- But 123.123 has **infinite digits** after transformed to binary

A Real Study V

- Compiler options can help to make things more consistent. For example, `-ffloat-store`: “Do not store floating-point variables in registers, and inhibit other options that might change whether a floating-point value is taken from a register or memory.”

```
$gcc -ffloat-store test.c; ./a.out
```

```
123.1230010986
```

```
15159.2734375000
```

```
123.1250000000
```

```
15159.7656250000
```

```
$gcc -ffloat-store -m32 test.c; ./a.out
```

```
123.1230010986
```

A Real Study VI

15159.2734375000

123.1250000000

15159.7656250000

- Note that other issues such as order of operations can also affect results.
- Consider running a real example using a machine learning software LIBSVM
- 64 bit:

A Real Study VII

```
$ ./svm-train -c 100 -e 0.00001 heart_scale
.....*..*
optimization finished, #iter = 2872
nu = 0.148045
obj = -2526.925470, rho = 1.145512
nSV = 107, nBSV = 9
Total nSV = 107
```

- 32bit:

A Real Study VIII

```
$ ./svm-train -c 100 -e 0.00001 heart_scale
.....*.*
optimization finished, #iter = 2819
nu = 0.148045
obj = -2526.925470, rho = 1.145515
nSV = 107, nBSV = 9
Total nSV = 107
```

- They are **different**
- Adding `-ffloat-store -mfpmath=387` is **not enough**
- 64 bit:

A Real Study IX

```
$ make clean; make
$ ./svm-train -c 100 -e 0.00001 heart_scale
rm -f *~ svm.o svm-train svm-predict svm-scale
g++ -Wall -Wconversion -O3 -fPIC -ffloat-store
g++ -Wall -Wconversion -O3 -fPIC -ffloat-store
g++ -Wall -Wconversion -O3 -fPIC -ffloat-store
g++ -Wall -Wconversion -O3 -fPIC -ffloat-store
.....*..*
optimization finished, #iter = 2863
nu = 0.148045
obj = -2526.925470, rho = 1.145512
nSV = 107, nBSV = 9
```

A Real Study X

Total nSV = 107

- We also need to **disable all optimization**
- 64bit:

```
$ make clean; make
```

```
$ ./svm-train -c 100 -e 0.00001 heart_scale
```

```
rm -f *~ svm.o svm-train svm-predict svm-scale
```

```
g++ -ffloat-store -mfpmath=387 -c svm.cpp
```

```
g++ -ffloat-store -mfpmath=387 svm-train.c s
```

```
g++ -ffloat-store -mfpmath=387 svm-predict.c
```

```
g++ -ffloat-store -mfpmath=387 svm-scale.c -
```

```
.....*...*
```

A Real Study XI

```
optimization finished, #iter = 3051
nu = 0.148045
obj = -2526.925470, rho = 1.145515
nSV = 107, nBSV = 9
Total nSV = 107
```

- 32 bit:

```
$ make clean; make
$ ./svm-train -c 100 -e 0.00001 heart_scale
rm -f *~ svm.o svm-train svm-predict svm-scale
g++ -m32 -ffloat-store -mfpmath=387 -c svm.o
g++ -m32 -ffloat-store -mfpmath=387 svm-train
```

A Real Study XII

```
g++ -m32 -ffloat-store -mfpmath=387 svm-prec  
g++ -m32 -ffloat-store -mfpmath=387 svm-scal  
.....*...*  
optimization finished, #iter = 3051  
nu = 0.148045  
obj = -2526.925470, rho = 1.145515  
nSV = 107, nBSV = 9  
Total nSV = 107
```

l